

rapport

October 22, 2020

1 Report 1 : Regression and Classification

Clément Antheaume, Camille-Amaury Juge.

```
[3]: from IPython.display import Image  
display(Image(filename='nantes.png'))
```



UNIVERSITÉ DE NANTES

1.1 Definitions of indicators :

1.1.1 Classification (categorical class)

Considering TP = Numbers of individuals well classified positive, TN = Numbers of individuals well classified negative and FP Numbers of positive individuals classified negative, FN Numbers of negative individuals classified positive.

Precision

- $TP / TP + TN$

Recall

- $TP / TP + FN$
- Represent the sensitivity of the model.

Specificity

- $TN / FP + TN$

Sensitivity

- $TP / TP + FN$

F-measure

- $2 * (Precision * Recall) / (Precision + Recall)$
- Mathematically, the harmonic mean of recall and precision.

Rand index

- $TP + TN / TP + FP + FN + TN$
- Percentage of correct decisions made by the classification algorithm.
- Can be used in clustering to measure the similarity between two clusters.

ROC Curve

- Function giving the number of True positive rate (y) given the false negative rate.
- The goal is to have a curve as close as possible to $y = x$.

1.1.2 Regression (numeric class)

Mean Squared Error

- MSE belongs to R^+
- $MSE = Average(Indicators - Indicator^2)$
- $\Leftrightarrow MSE = Bias(Indicator)^2 + Variance(Indicator)$
- As we can see, it can be defined as a measure of the bias and variance of the Indicator
- It evaluates the quadratic risk of the Indicator
- Sensitive to outliers (large error values), thus useful when we want our model to be quite stable

Root Mean Squared-Error

- RMSE belongs to R^+
- $RMSE = Root(MSE) = Root(Average(Indicators - Indicator^2))$
- $\Leftrightarrow RMSE = Root(Bias(Indicator)^2 + Variance(Indicator))$
- As we can see, it can be defined as a measure of the standard deviation of the Indicator
- It evaluates the quadratic risk of the Indicator
- Even more Sensitive to outliers (large error values), thus useful when we want our model to avoid large errors

Mean Bias Error

- MBE belongs to \mathbb{R}
- $MBE = \text{Average}(Y_{\text{label}} - Y_{\text{predicted}})$
- As we can see, it can be defined as a measure of the bias of the error between labels and predictions
- It indicates if the model overestimates (if $MBE < 0$) or underestimates (if $MBE > 0$) the output

Systematic Error

- SE or SD belongs to \mathbb{R}^+
- $SD = \sqrt{\text{RMSE}(\text{error})^2 - \text{MBE}(\text{error})^2}$
- As we can see, it can be defined as a measure of the MSE-Bias so it reduces the importance of larger errors

Mean Absolute Error

- MAE belongs to \mathbb{R}^+
- $MAE = \text{Average}(|Y_{\text{label}} - Y_{\text{predicted}}|)$
- As we can see, it can be defined as a measure of the bias, not regarding to its orientation
- MAE will not be sensible to outliers

Mean Absolute Percentage Error

- MAPE belongs to \mathbb{R}^+
- $MAPE = \text{Average}(|Y_{\text{label}} - Y_{\text{predicted}}| / |Y_{\text{label}}|)$
- As we can see, it can be defined as a measure of the bias, not regarding to its orientation
- It has the advantage to show ratio errors rather than value errors

R2

- R^2 belongs to \mathbb{R} , R^2 belongs to $[-1, 1]$
- $R^2 = \text{Correlation}(Y_{\text{predicted}}, Y_{\text{label}})$
- $\Leftrightarrow R^2 = \frac{\text{Sum}((Y_{\text{predicted}} - \text{Average}(Y_{\text{label}}))^2)}{\text{Sum}((Y_{\text{label}} - \text{Average}(Y_{\text{label}}))^2)}$
- As we can see, it can be defined as a measure of the correlation of the error
- It has the advantages to put every error on the same scale

1.1.3 Validation Techniques

Hold Out Cross Validation Separate the dataset in two sub-datasets : * Training Set is used to train the model. * Testing Set is used to validate the model with indicators.

The splitting is done with a percentage of the initial dataset (for instance 80%/20%).

This method has the advantage to avoid overfitting. But it is not stable since we still have a low probability to have the worst configuration in our sub-datasets.

K-Fold Cross Validation Separate the dataset S in k sub-datasets, then each subset contains N/k individuals. Thus, we iterate k times on all sub-datasets : * we create a training set with $k-1$ sub-datasets * we create a testing set with the last sub-dataset * we compute the empirical error

When the k-iterations are done, we compute the mean of the empirical error. Then we have a stable validation indicator over multiple configuration (k) of our datasets.

1.2 Classification with Python

1.2.1 Imports

```
[100]: import pandas as pd
import numpy as np
from sklearn import tree
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn import datasets
import random
import disarray
```

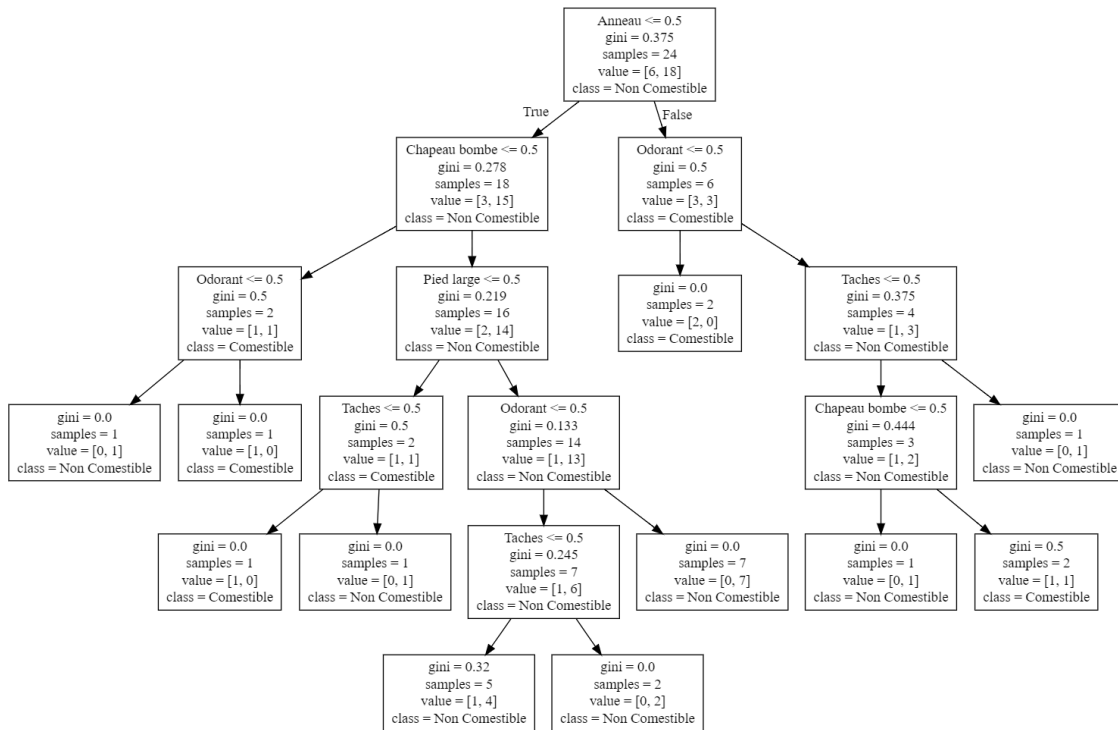
1.2.2 Mushroom Decision Tree

I modified the Mushroom dataset to be a xlsx because there were way too many different inputs and I couldn't find one that would make the process easier to us. Here's the download link for the xls version : <https://tinyurl.com/y68krwbw>

```
[31]: classes = [u'Comestible', u'Non Comestible']
data = pd.read_csv('Mushroom.csv', encoding='utf-8')
X = data.iloc[:, 1:-1]
Y = data.iloc[:, -1]
dtree = tree.DecisionTreeClassifier()
dtree = dtree.fit(X, Y)

[70]: tree.export_graphviz(dtree, out_file="mushroom.dot", feature_names=X.columns,
    →class_names=classes)

[71]: display(Image(filename='mushroom.png'))
```



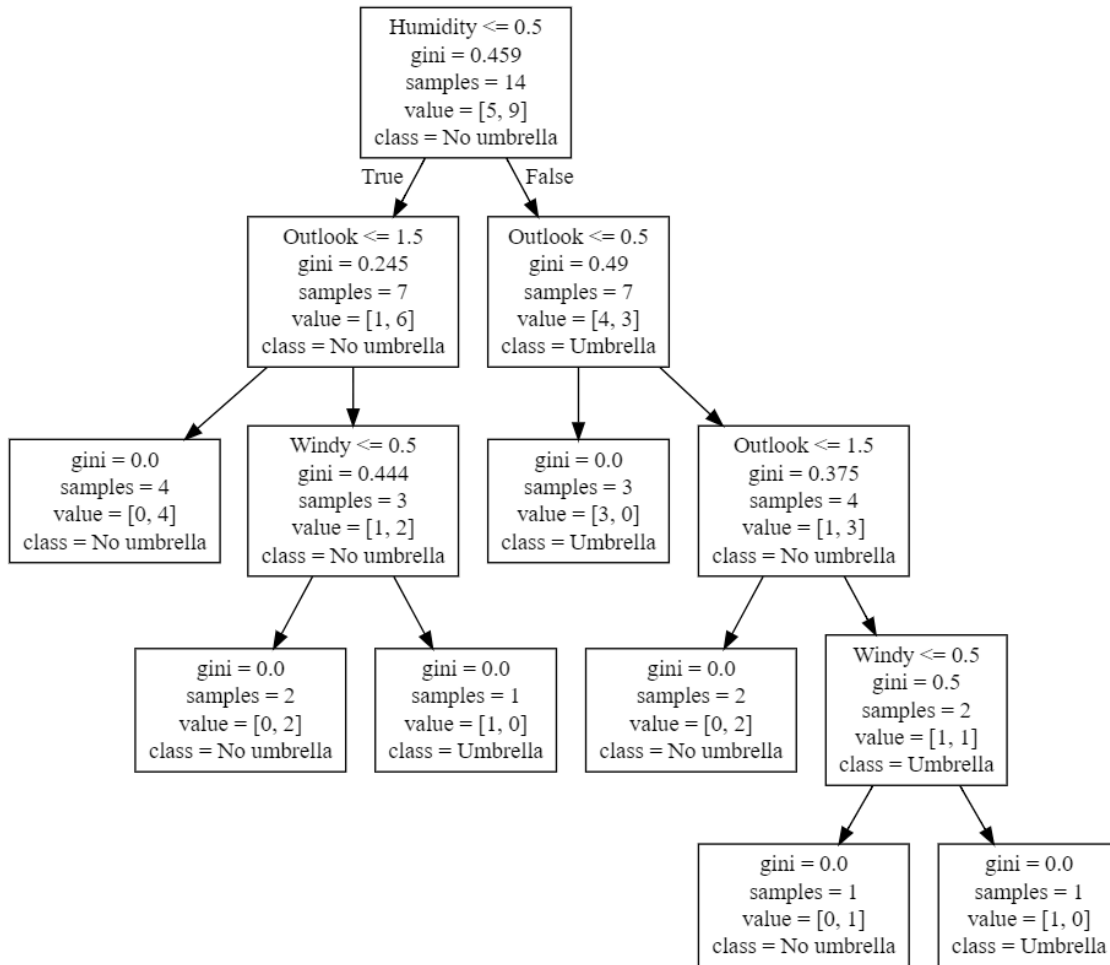
1.2.3 Weather Decision Tree

```
[112]: classes = ['Umbrella', 'No umbrella']
data = pd.read_excel('Meteo.xls').replace({True : 1, False : 0, 'high' : 1, 'normal' : 0, 'N' : 0, 'P' : 1})
data = pd.get_dummies(data, columns=["Temperature", "Outlook"])
X = data.iloc[:,1:-1]
Y = data.iloc[:, -1]

dtree = tree.DecisionTreeClassifier()
dtree = dtree.fit(X, Y)
```

```
[117]: tree.export_graphviz(dtree, out_file="meteo.dot", feature_names=X.columns,
    class_names=classes)
```

```
[118]: display(Image(filename='meteo.png'))
```



1.2.4 Some comparison on the Iris dataset

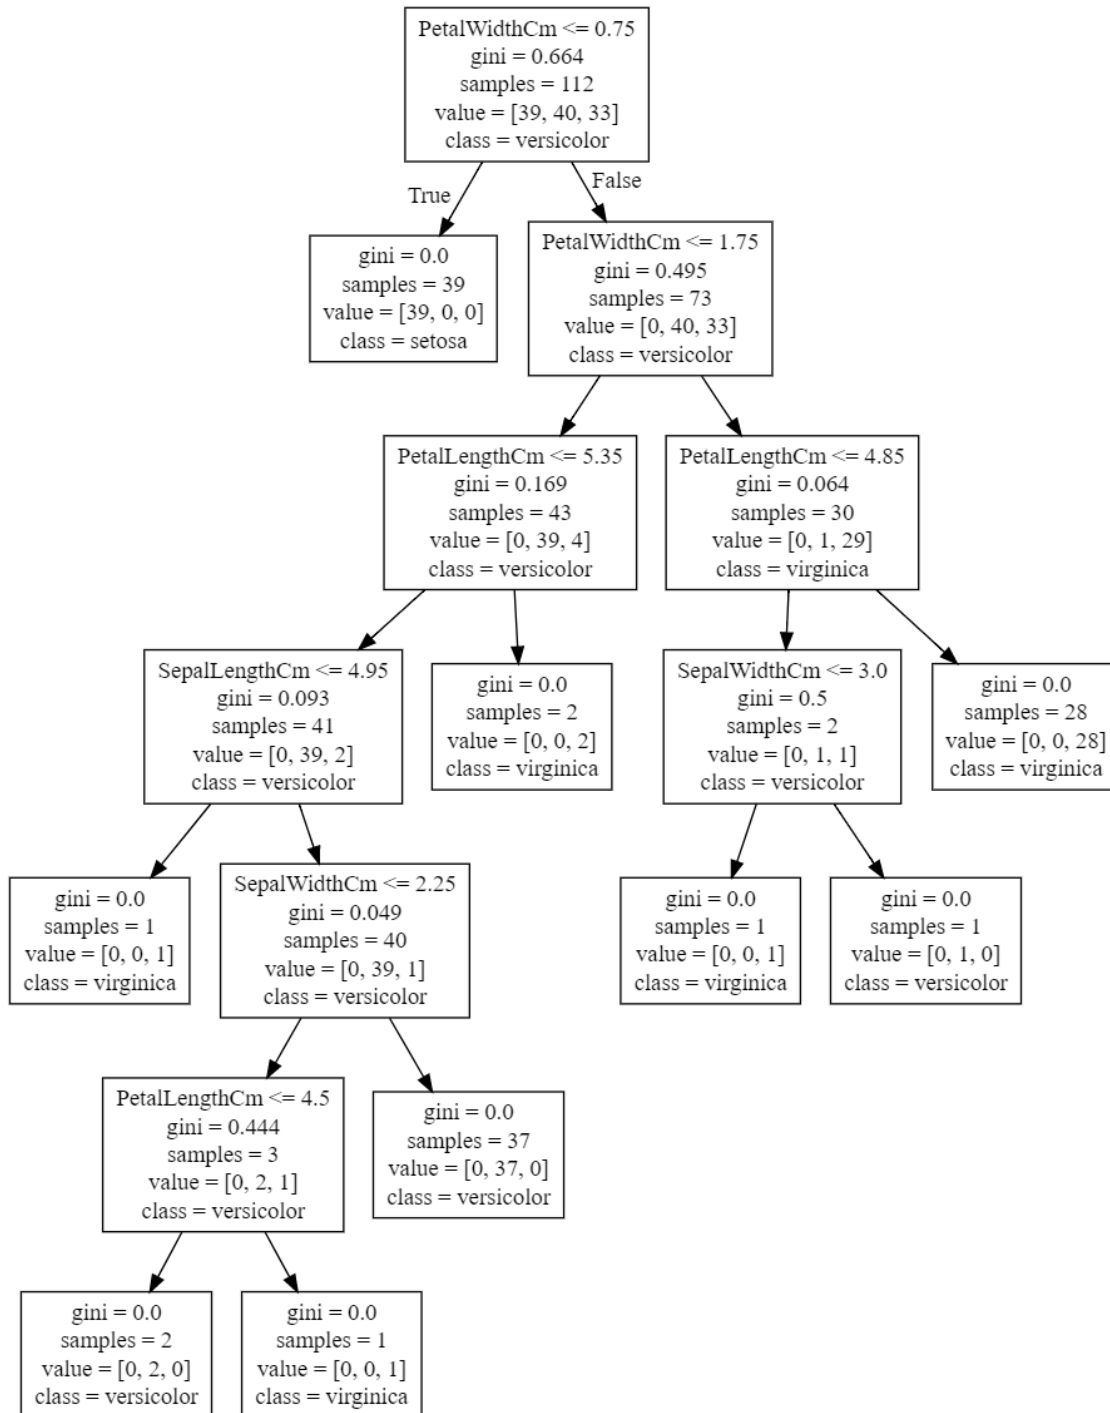
Decision Tree Classification

```
[151]: iris = datasets.load_iris()
X = pd.DataFrame(iris.data, columns = iris.feature_names)
Y = pd.DataFrame(iris.target)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y)
classes = iris.target_names
```

```
[152]: dtree = tree.DecisionTreeClassifier()
dtree = dtree.fit(X_train, Y_train)
```

```
[153]: tree.export_graphviz(dtree, out_file="iris.dot", feature_names=features,
    ↳class_names=classes)
```

```
[154]: display(Image(filename='iris.png'))
```



```

[155]: def get_metrics(conf): # Compute the generic metrics from the confusion matrix
    # from each class in the matrix
    # Computing indicators sub-terms
    return pd.DataFrame(conf).da.export_metrics()
  
```

```
[156]: Y_pred = dtree.predict(X_test)
conf_ = confusion_matrix(Y_test, Y_pred).astype(int)
get_metrics(conf_)
```

```
[156]:
```

	0	1	2	micro-average
accuracy	1.0	0.868421	0.868421	0.912281
f1	1.0	0.827586	0.705882	0.868421
false_discovery_rate	0.0	0.250000	0.142857	0.131579
false_negative_rate	0.0	0.076923	0.400000	0.131579
false_positive_rate	0.0	0.160000	0.035714	0.065789
negative_predictive_value	1.0	0.954545	0.870968	0.934211
positive_predictive_value	1.0	0.750000	0.857143	0.868421
precision	1.0	0.750000	0.857143	0.868421
recall	1.0	0.923077	0.600000	0.868421
sensitivity	1.0	0.923077	0.600000	0.868421
specificity	1.0	0.840000	0.964286	0.934211
true_negative_rate	1.0	0.840000	0.964286	0.934211
true_positive_rate	1.0	0.923077	0.600000	0.868421

We can see that the first class is clearly well predicted but both second and third class are 0.912...

It means that the model is able to correctly classify 947 out of 1000 times (rounded values)

Random Forest Classification

```
[135]: rf = AdaBoostClassifier(n_estimators=100, random_state=0)
rf.fit(X_train, np.ravel(Y_train))
```

```
[135]: AdaBoostClassifier(n_estimators=100, random_state=0)
```

```
[136]: Y_pred = rf.predict(X_test)

rf_conf = confusion_matrix(Y_test, Y_pred).astype(int)
get_metrics(rf_conf)
```

```
[136]:
```

	0	1	2	micro-average
accuracy	1.0	0.947368	0.947368	0.964912
f1	1.0	0.933333	0.909091	0.947368
false_discovery_rate	0.0	0.125000	0.000000	0.052632
false_negative_rate	0.0	0.000000	0.166667	0.052632
false_positive_rate	0.0	0.083333	0.000000	0.026316
negative_predictive_value	1.0	1.000000	0.928571	0.973684
positive_predictive_value	1.0	0.875000	1.000000	0.947368
precision	1.0	0.875000	1.000000	0.947368
recall	1.0	1.000000	0.833333	0.947368
sensitivity	1.0	1.000000	0.833333	0.947368
specificity	1.0	0.916667	1.000000	0.973684
true_negative_rate	1.0	0.916667	1.000000	0.973684
true_positive_rate	1.0	1.000000	0.833333	0.947368

We can see that the first class is clearly well predicted but both second and third class are 0.947...

It means that the model is able to correctly classify 964 out of 1000 times (rounded values)

The prediction is then better than the decision tree.

AdaBoost Classification

```
[137]: ada = AdaBoostClassifier(n_estimators=100, random_state=0)
ada.fit(X_train, np.ravel(Y_train))
```

```
[137]: AdaBoostClassifier(n_estimators=100, random_state=0)
```

```
[138]: Y_pred = ada.predict(X_test)

ada_conf = confusion_matrix(Y_test, Y_pred).astype(int)
get_metrics(ada_conf)
```

```
[138]:
```

	0	1	2	micro-average
accuracy	1.0	0.947368	0.947368	0.964912
f1	1.0	0.933333	0.909091	0.947368
false_discovery_rate	0.0	0.125000	0.000000	0.052632
false_negative_rate	0.0	0.000000	0.166667	0.052632
false_positive_rate	0.0	0.083333	0.000000	0.026316
negative_predictive_value	1.0	1.000000	0.928571	0.973684
positive_predictive_value	1.0	0.875000	1.000000	0.947368
precision	1.0	0.875000	1.000000	0.947368
recall	1.0	1.000000	0.833333	0.947368
sensitivity	1.0	1.000000	0.833333	0.947368
specificity	1.0	0.916667	1.000000	0.973684
true_negative_rate	1.0	0.916667	1.000000	0.973684
true_positive_rate	1.0	1.000000	0.833333	0.947368

We can see that the first class is clearly well predicted but both second and third class are 0.947...

It means that the model is able to correctly classify 964 out of 1000 times (rounded values)

The prediction is then better than the decision tree.

1.3 Linear Regression with R

The dataset : <https://www.kaggle.com/mohansacharya/graduate-admissions>

The Content :

The dataset contains several parameters which are considered important during the application for Masters Programs. The parameters included are :

- GRE Scores (out of 340)
- TOEFL Scores (out of 120)
- University Rating (out of 5)
- Statement of Purpose and Letter of Recommendation Strength (out of 5)
- Undergraduate GPA (out of 10)

- Research Experience (either 0 or 1)
- Chance of Admit (ranging from 0 to 1)

1.3.1 Imports

```
[1]: library(namespace)
registerNamespace('psy', loadNamespace('psych'))
library(ggplot2)
library(reshape2)
library(lattice)
registerNamespace('ml', loadNamespace('caret'))
registerNamespace('metrics', loadNamespace('Metrics'))
registerNamespace('mlmetrics', loadNamespace('MLmetrics'))
library("IRdisplay")
```

```
<environment: namespace:psych>
```

```
Registered S3 methods overwritten by 'ggplot2':
```

```
method      from
[.quosures   rlang
c.quosures   rlang
print.quosures rlang
```

```
<environment: namespace:caret>
```

```
<environment: namespace:Metrics>
```

```
<environment: namespace:MLmetrics>
```

```
[2]: csv <- read.csv("Admission_Predict.csv", header = TRUE)
head(csv[,2:ncol(csv)])
```

GRE.Score	TOEFL.Score	University.Rating	SOP	LOR	CGPA	Research	Chance.of.Admit
337	118	4	4.5	4.5	9.65	1	0.92
324	107	4	4.0	4.5	8.87	1	0.76
316	104	3	3.0	3.5	8.00	1	0.72
322	110	3	3.5	2.5	8.67	1	0.80
314	103	2	2.0	3.0	8.21	0	0.65
330	115	5	4.5	3.0	9.34	1	0.90

1.3.2 Data Understanding

Univariate Analysis

```
[3]: summ <- psy::describe(csv[,2:ncol(csv)])
summ[,1:(ncol(summ)%/%2 + 1)]
```


	vars	n	mean	sd	median	trimmed	mad
GRE.Score	1	400	316.807500	11.4736461	317.00	316.8500000	11.860800
TOEFL.Score	2	400	107.410000	6.0695138	107.00	107.3281250	5.930400
University.Rating	3	400	3.087500	1.1437281	3.00	3.0656250	1.482600
SOP	4	400	3.400000	1.0068686	3.50	3.4296875	0.741300
LOR	5	400	3.452500	0.8984775	3.50	3.4625000	0.741300
CGPA	6	400	8.598925	0.5963171	8.61	8.6024687	0.667170
Research	7	400	0.547500	0.4983620	1.00	0.5593750	0.000000
Chance.of.Admit	8	400	0.724350	0.1426093	0.73	0.7309688	0.133434

```
[4]: summ[, (ncol(summ)%%2 + 1):ncol(summ)]
```

	mad	min	max	range	skew	kurtosis	se
GRE.Score	11.860800	290.00	340.00	50.00	-0.06242254	-0.7181786	0.573682306
TOEFL.Score	5.930400	92.00	120.00	28.00	0.05678751	-0.5985838	0.303475689
University.Rating	1.482600	1.00	5.00	4.00	0.16997797	-0.8123104	0.057186406
SOP	0.741300	1.00	5.00	4.00	-0.27369641	-0.6937320	0.050343432
LOR	0.741300	1.00	5.00	4.00	-0.10619038	-0.6808341	0.044923877
CGPA	0.667170	6.80	9.92	3.12	-0.06549644	-0.4803728	0.029815855
Research	0.000000	0.00	1.00	1.00	-0.19014793	-1.9687469	0.024918099
Chance.of.Admit	0.133434	0.34	0.97	0.63	-0.35080166	-0.4122290	0.007130467

We can see that : * For each Variable except Research, the mean and median are quite the same : it means that there are no outliers which make mean varying much. Then, we conclude that since Research mean is near 0.54 and median 1., there are a lot of outliers values (near 0) which tend to attract mean. * The standard deviation is rather small for some variable comparing to their range of values, meaning that this variable has quite regrouped individuals values. * here is a quick interpretation of the kurtosis values :

```
[5]: display_png(file="kurtosis.png")
```




We can see that for example, the TOEFL.Score is near to follow a distribution Law of Cosinus (around -0.593762).

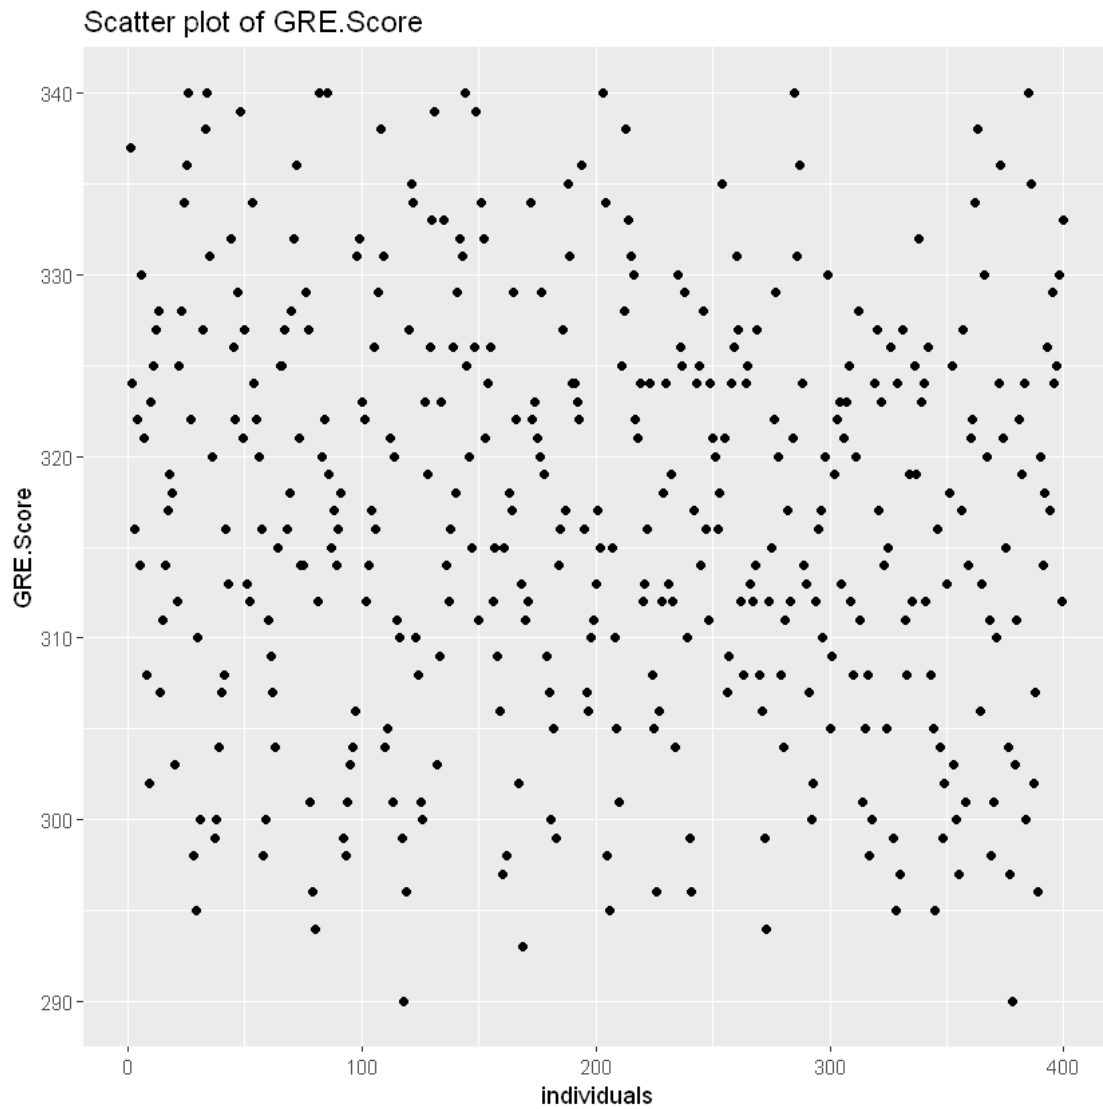
- The skew indicator helps us to know in which direction the 'tail' of the asymmetric or symmetric distribution is going. We will take the example of Research again, which confirms the fact that the low values represent the tail following the negative skew coefficient.

Let's check for outliers over our dataset :

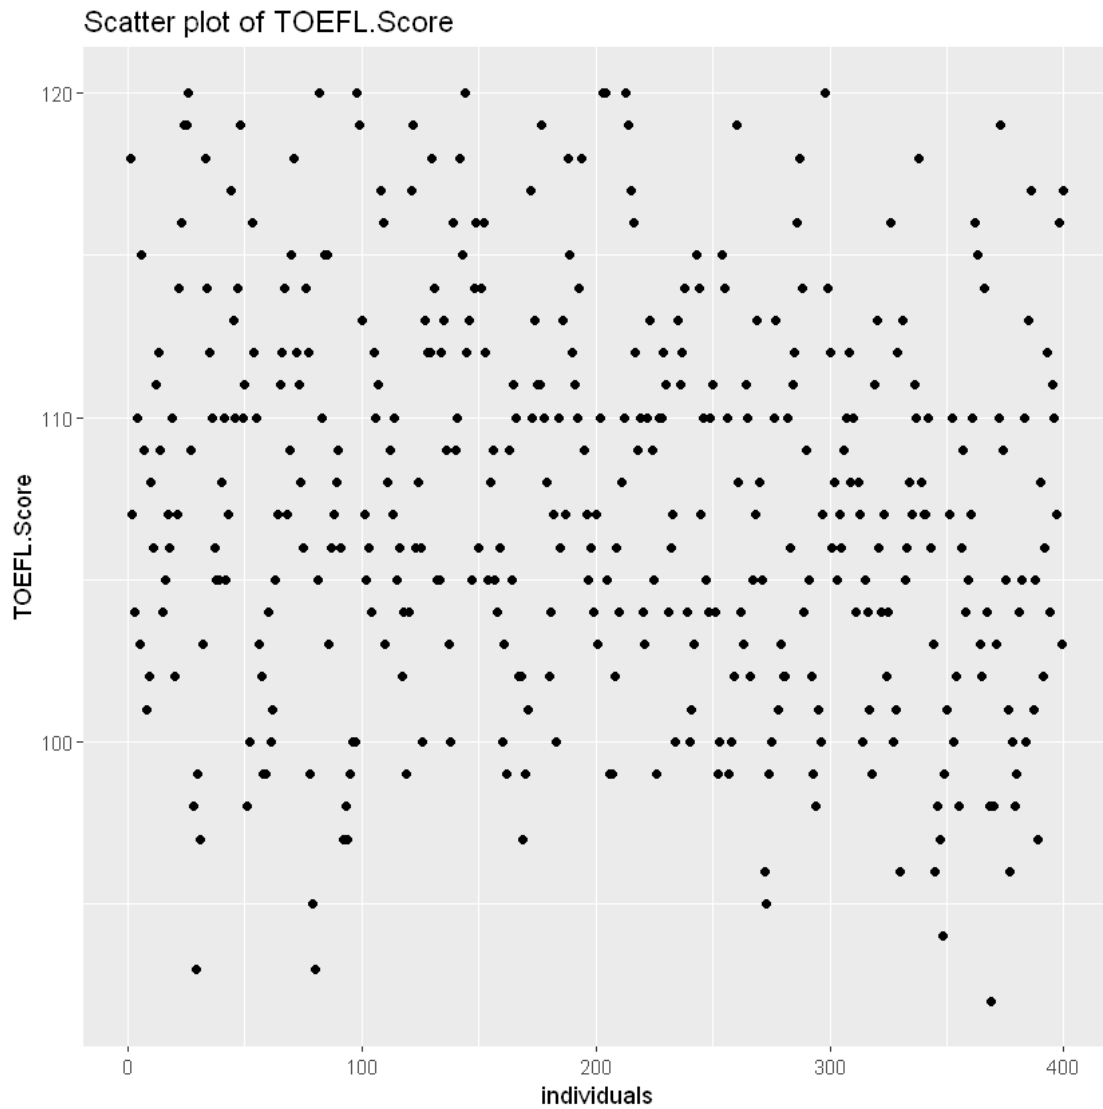
```
[7]: boxplot(csv)
```



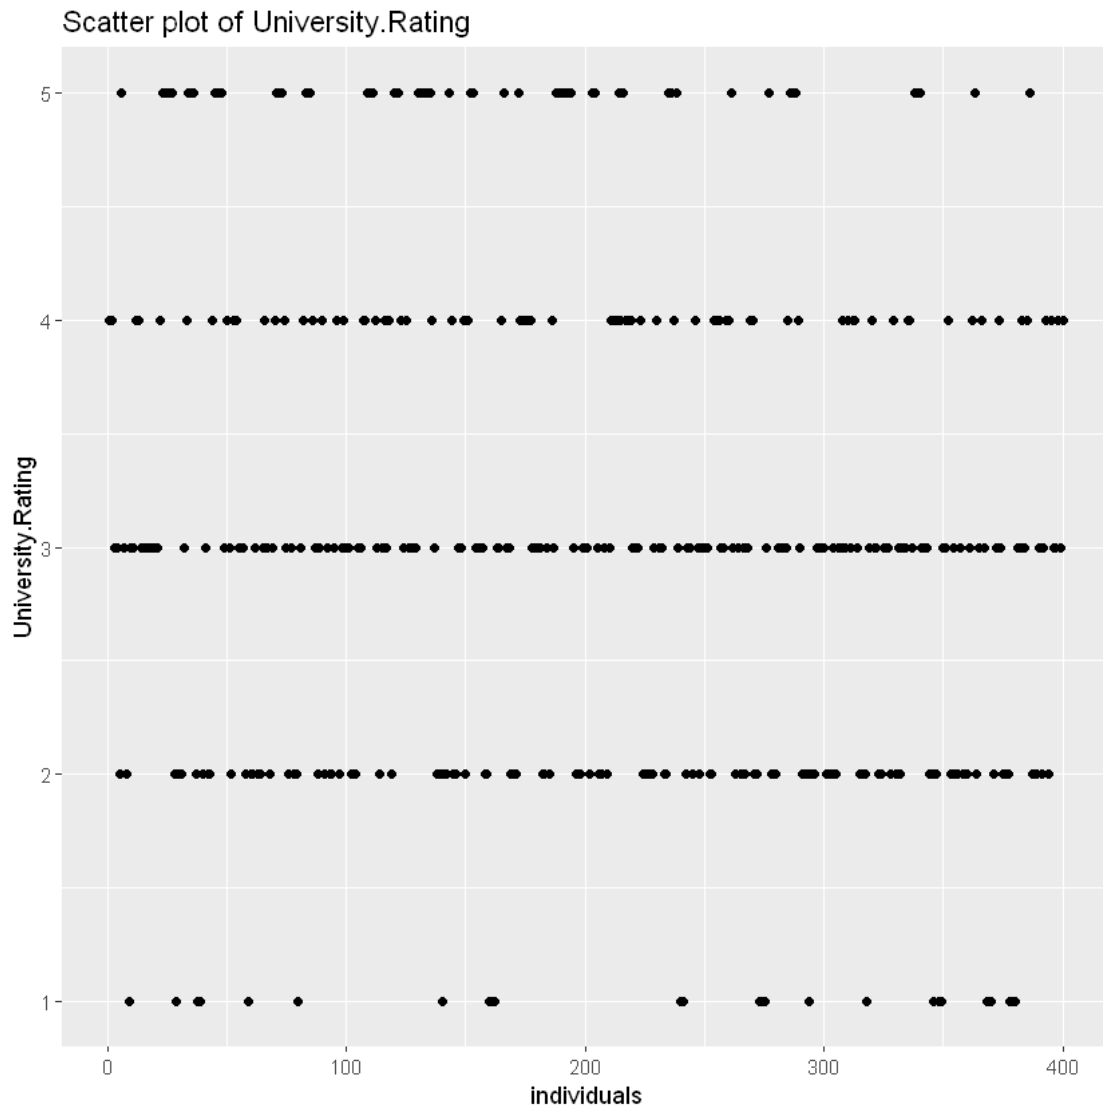
```
[8]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),2])) + geom_point() +
      labs(title = "Scatter plot of GRE.Score",
           x = "individuals", y = "GRE.Score")
```




```
[9]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),3])) + geom_point() +  
      labs(title = "Scatter plot of TOEFL.Score",  
           x = "individuals", y = "TOEFL.Score")
```



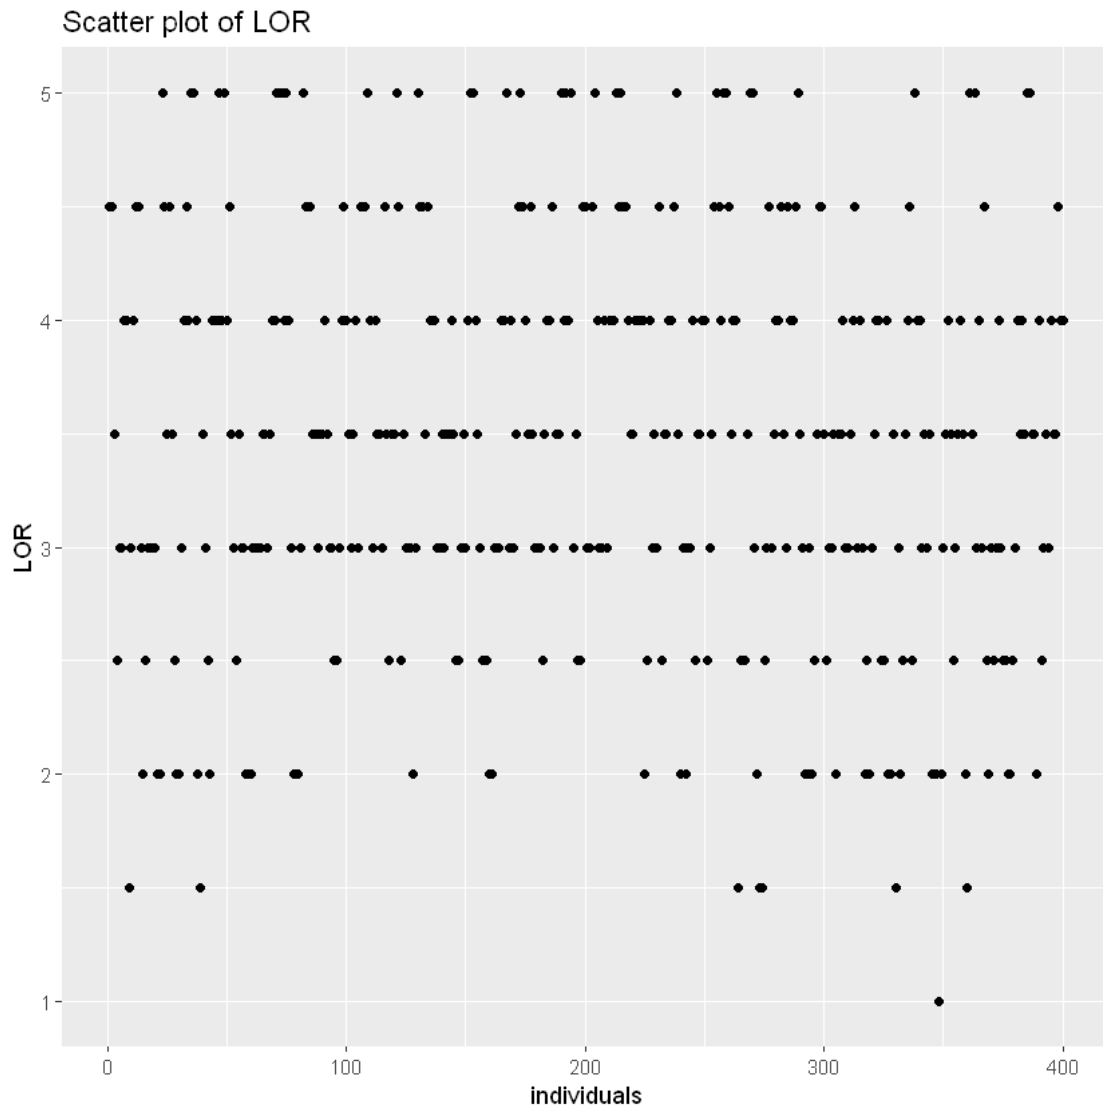
```
[10]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),4])) + geom_point() +
      labs(title = "Scatter plot of University.Rating",
           x = "individuals", y = "University.Rating")
```



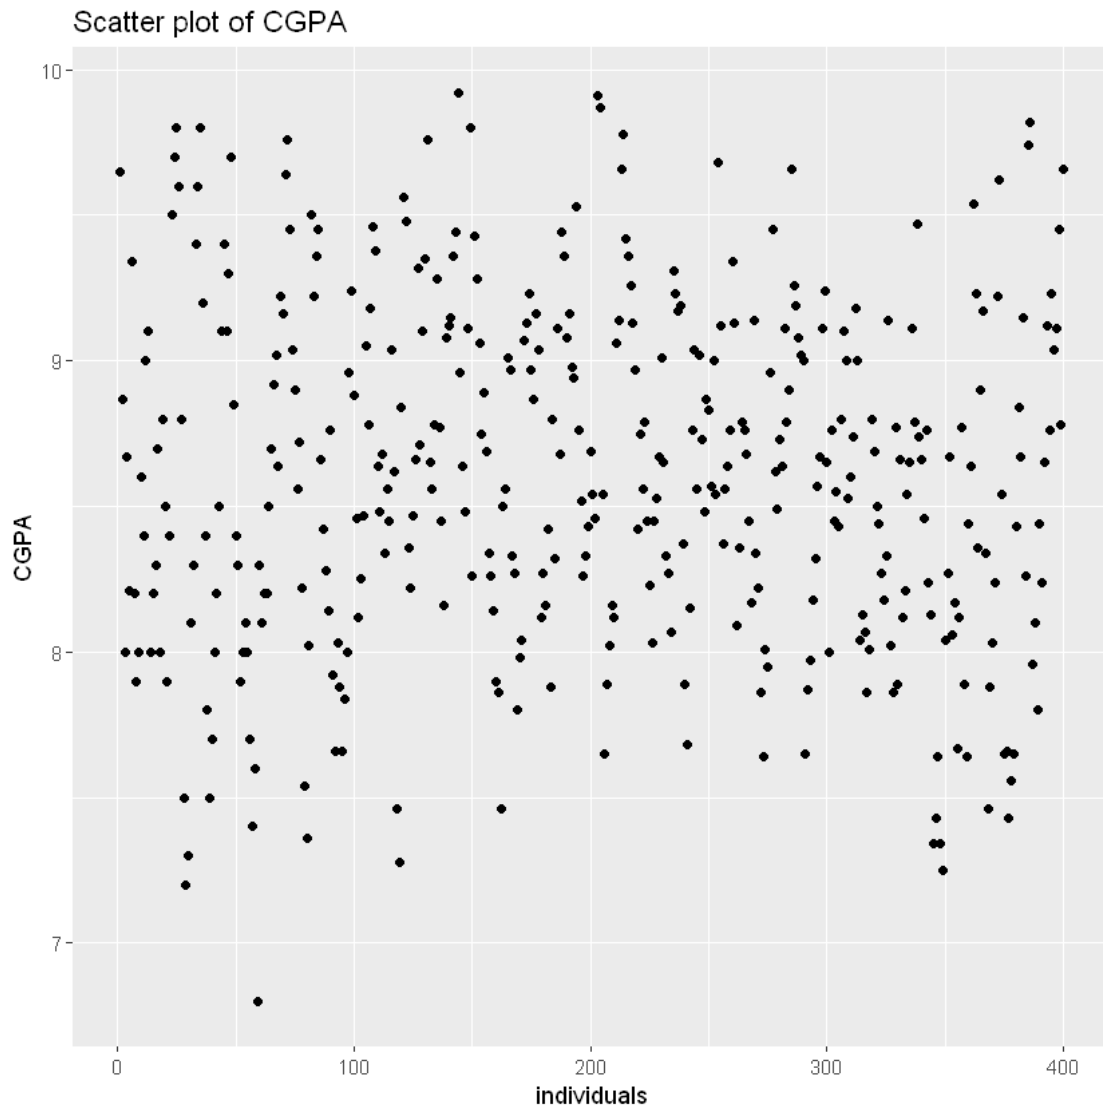
```
[11]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),5])) + geom_point() +
      labs(title = "Scatter plot of SOP", x = "individuals", y = "SOP")
```

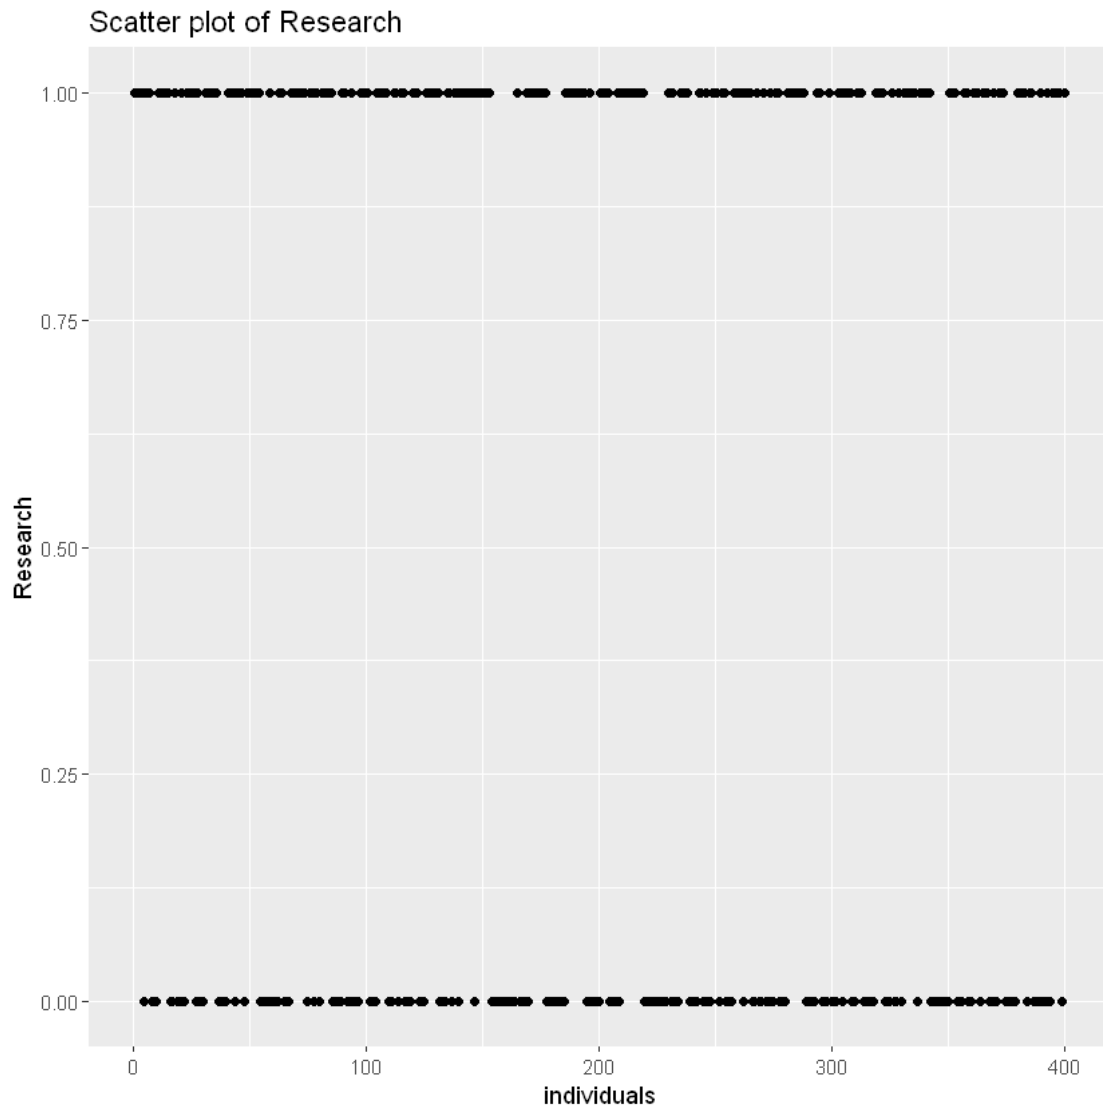
```
[12]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),6])) + geom_point() +
      labs(title = "Scatter plot of LOR", x = "individuals", y = "LOR")
```



```
[13]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),7])) + geom_point() +
      labs(title = "Scatter plot of CGPA", x = "individuals", y = "CGPA")
```



```
[14]: ggplot(csv, aes(x = csv[1:nrow(csv),1], y = csv[1:nrow(csv),8])) + geom_point() +  
      labs(title = "Scatter plot of Research", x = "individuals", y = "Research")
```




We can see there are two methods to rank students : * The first one deals with continuous values, ranking people with numerical data. * The second one deals with categorical data (and more precisely ordinal data), using discrete action space.


Those plots do not really helps us to interpret more informations but rather confirms what we could have seen in the data describing.

```
[15]: correlation_pearson_csv <- melt(cor(csv[1:nrow(csv),0:-1],
                                     method = c("pearson")))
correlation_kendall_csv <- melt(cor(csv[1:nrow(csv),0:-1],
                                     method = c("kendall")))
correlation_spearman_csv <- melt(cor(csv[1:nrow(csv),0:-1],
                                     method = c("spearman")))
```

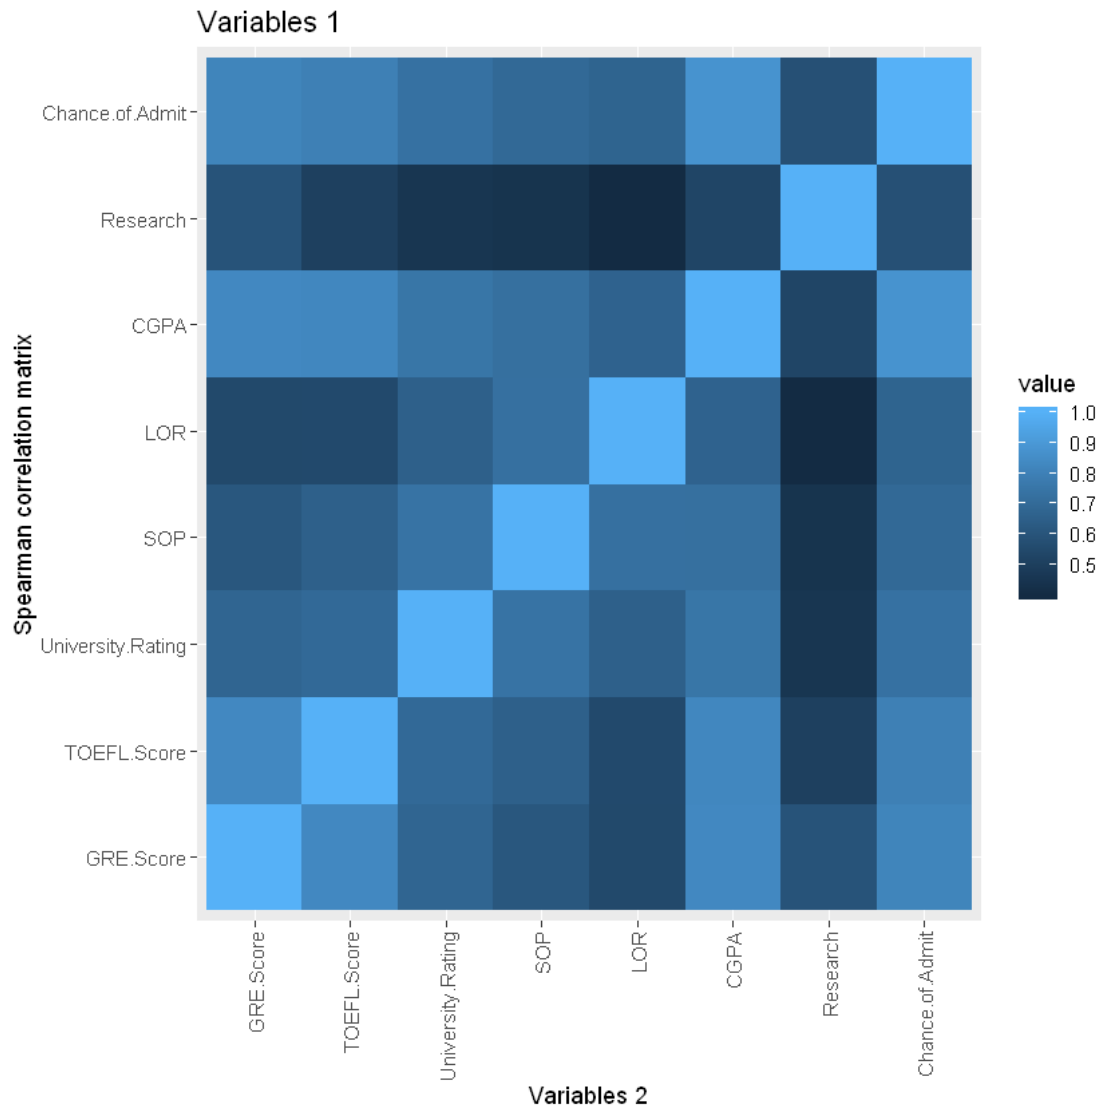
```
[16]: ggplot(data = correlation_pearson_csv, aes(x=Var1, y=Var2, fill=value)) +  
  geom_tile() +  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1)) +  
  labs(title = "Variables 1", x = "Variables 2",  
        y = "Pearson correlation matrix")
```



```
[17]: ggplot(data = correlation_kendall_csv, aes(x=Var1, y=Var2, fill=value)) +  
  geom_tile() +  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1)) +  
  labs(title = "Variables 1", x = "Variables 2",  
        y = "Kendall correlation matrix")
```



```
[18]: ggplot(data = correlation_spearman_csv, aes(x=Var1, y=Var2, fill=value)) +  
  geom_tile() +  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1)) +  
  labs(title = "Variables 1", x = "Variables 2",  
       y = "Spearman correlation matrix")
```

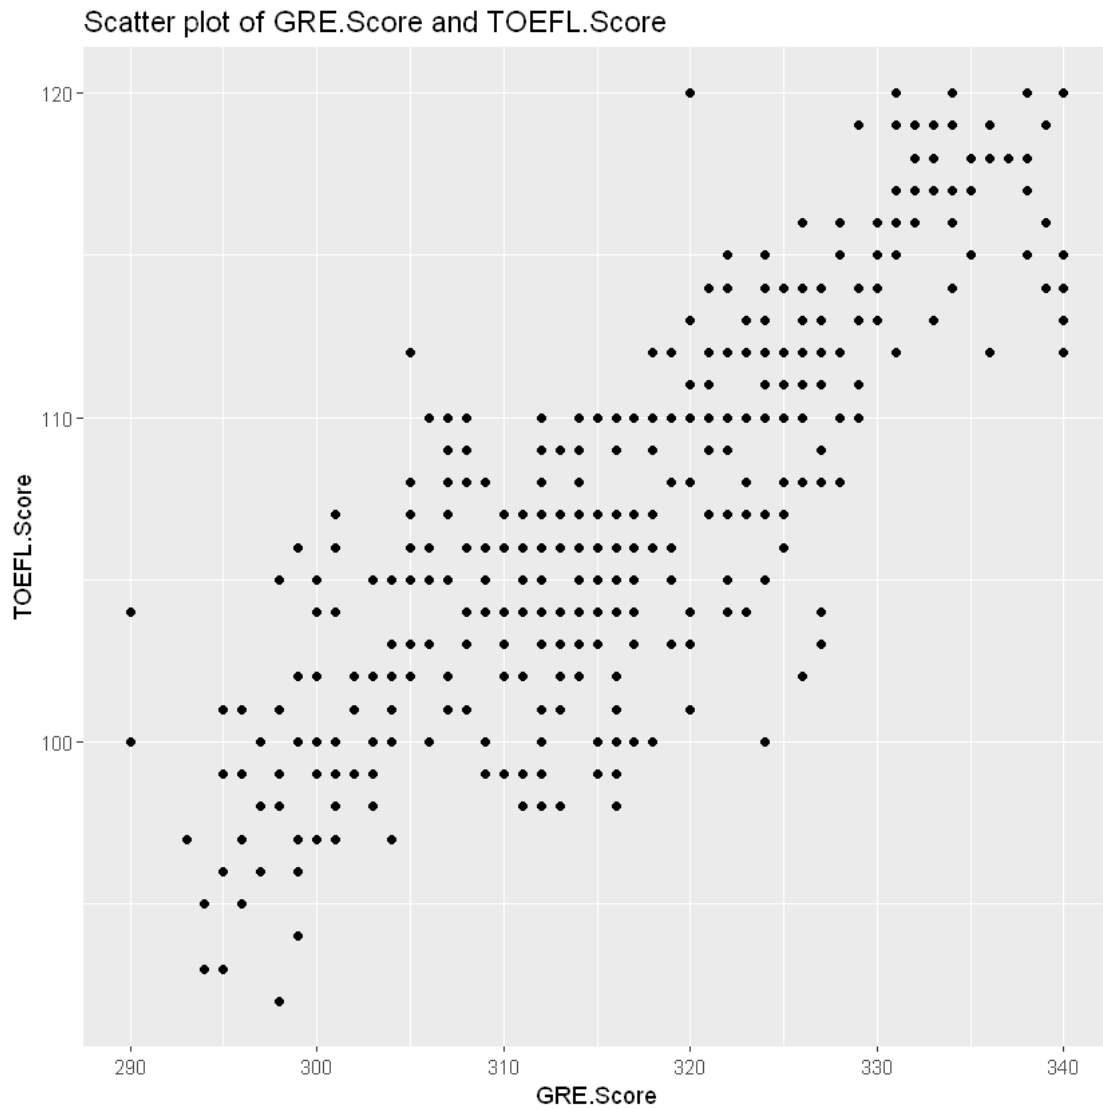


As we can quickly see, there are some linear correlation following the 3 indicators 'Pearson, Kendall and Spearman'.

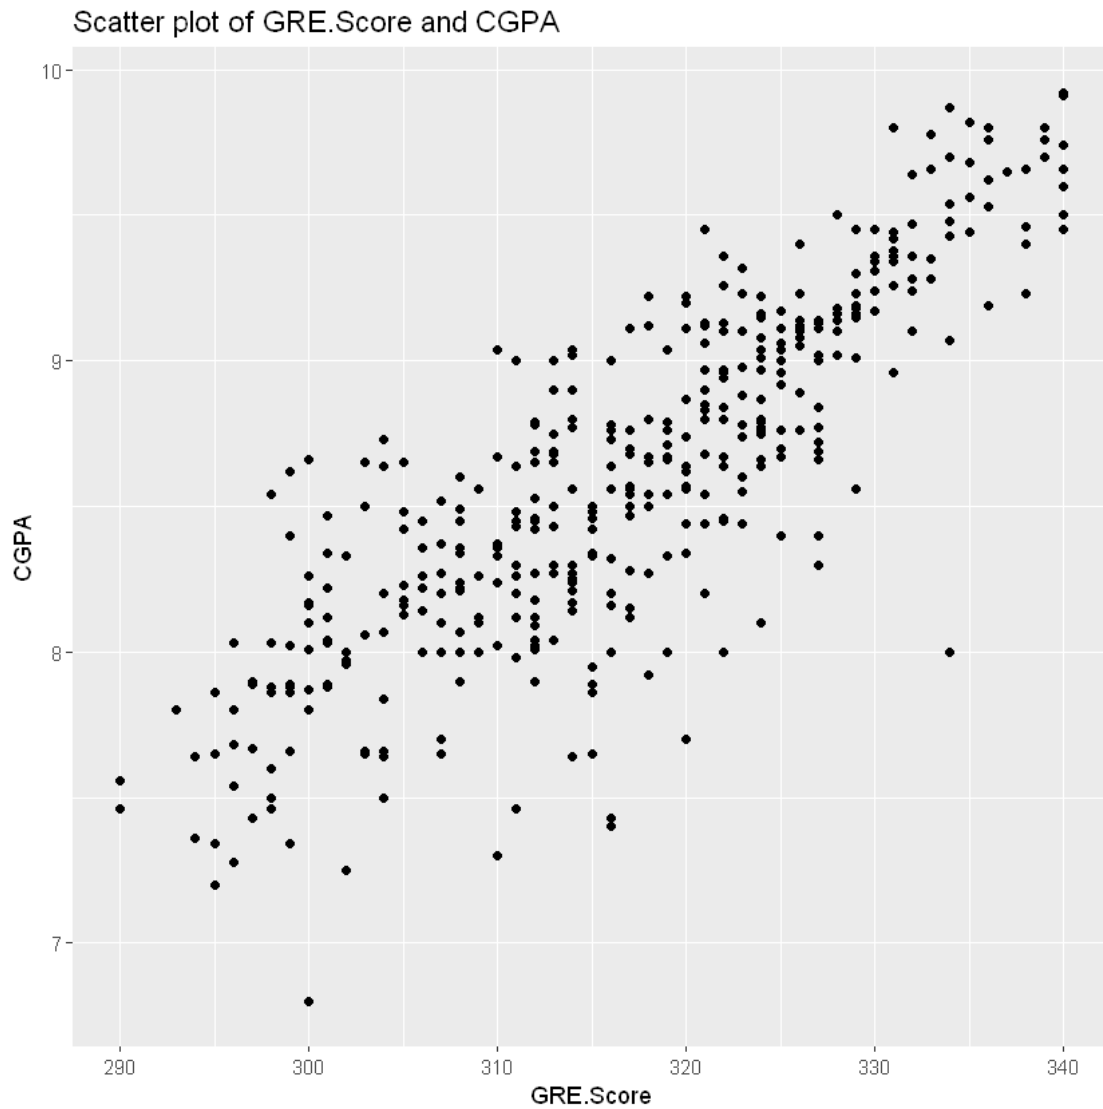
- GRE.Score, TOEFL.Score and CPGA are highly correlated one by one but also to Chance of Admits. This means that those scores highly determines the chances to be admitted by themselves.

Let's plot those variables in order to have a confirmation

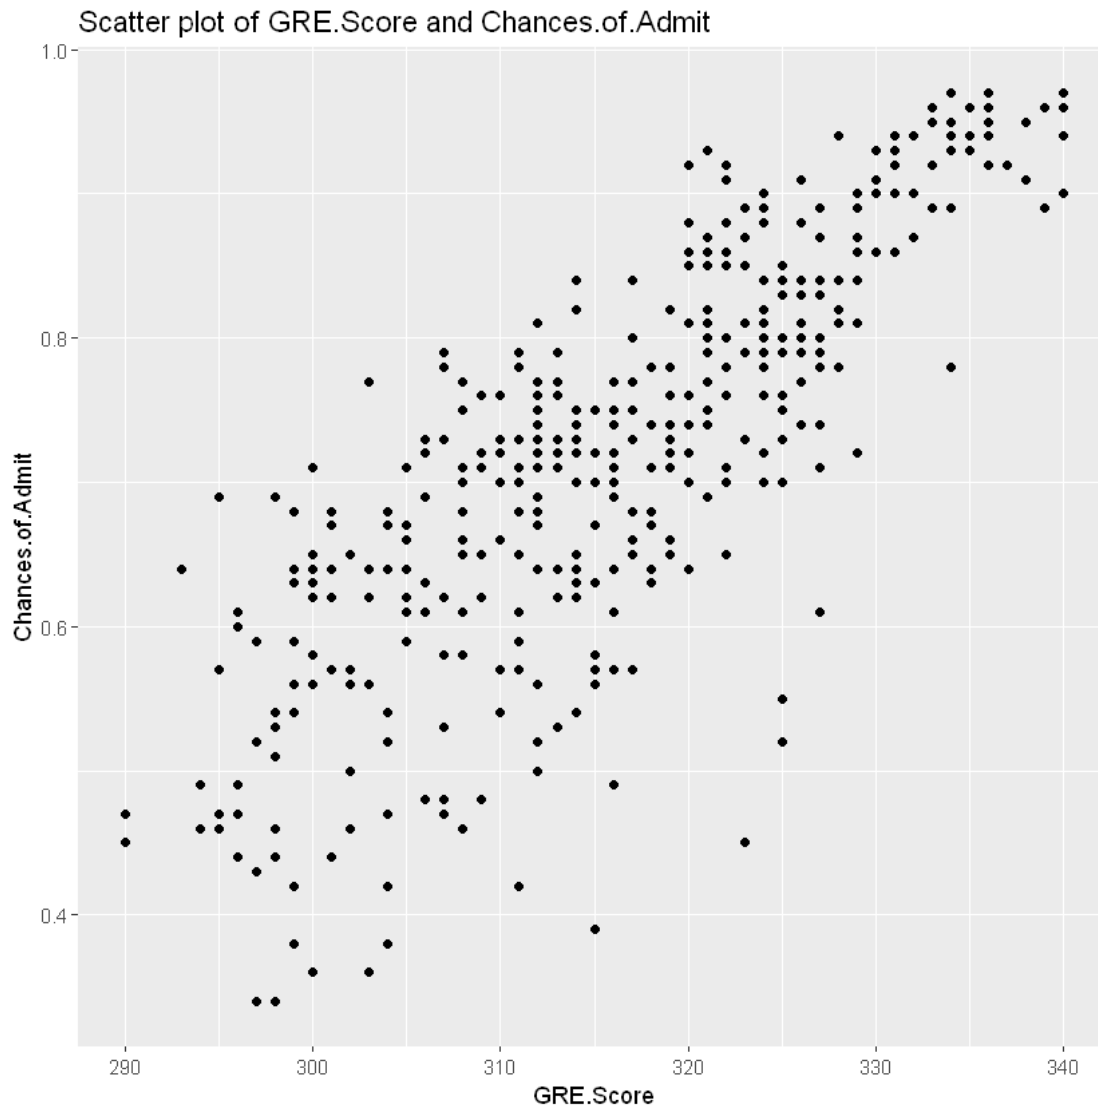
```
[19]: ggplot(csv, aes(x = csv[1:nrow(csv),2], y = csv[1:nrow(csv),3])) + geom_point() +
      labs(title = "Scatter plot of GRE.Score and TOEFL.Score",
           x = "GRE.Score", y = "TOEFL.Score")
```



```
[20]: ggplot(csv, aes(x = csv[1:nrow(csv),2], y = csv[1:nrow(csv),7])) + geom_point() +  
      labs(title = "Scatter plot of GRE.Score and CGPA", x = "GRE.Score",  
           y = "CGPA")
```

```
[21]: ggplot(csv, aes(x = csv[1:nrow(csv),2], y = csv[1:nrow(csv),9])) + geom_point() +  
      labs(title = "Scatter plot of GRE.Score and Chances.of.Admit",  
           x = "GRE.Score", y = "Chances.of.Admit")
```



We easily understand why linear regression could be a good approximation of this dataset seeing those last plots.

We could have continued with PCA and T-SNE data reduction in order to have more comprehension over the dataset. Since it's not the aim of this exercise, we will pursue with the linear regression.

1.3.3 Machine Learning

Re-arranging Data

```
[22]: # Data are numbers ?  
str(csv)  
# train/test splitting  
inTrain = ml::createDataPartition(y = csv[1:nrow(csv),9], p = .80,
```

```

                                list = FALSE)
train_csv <- csv[inTrain,2:ncol(csv)]
test_csv <- csv[-inTrain,2:ncol(csv)]
head(train_csv)
head(test_csv)
print(paste("Train : ",nrow(train_csv)))
print(paste("Test : ",nrow(test_csv)))

```

```

'data.frame':  400 obs. of  9 variables:
 $ Serial.No.      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ GRE.Score       : int  337 324 316 322 314 330 321 308 302 323 ...
 $ TOEFL.Score     : int  118 107 104 110 103 115 109 101 102 108 ...
 $ University.Rating: int  4 4 3 3 2 5 3 2 1 3 ...
 $ SOP            : num  4.5 4 3 3.5 2 4.5 3 3 2 3.5 ...
 $ LOR            : num  4.5 4.5 3.5 2.5 3 3 4 4 1.5 3 ...
 $ CGPA           : num  9.65 8.87 8 8.67 8.21 9.34 8.2 7.9 8 8.6 ...
 $ Research       : int  1 1 1 1 0 1 1 0 0 0 ...
 $ Chance.of.Admit : num  0.92 0.76 0.72 0.8 0.65 0.9 0.75 0.68 0.5 0.45 ...

```

	GRE.Score	TOEFL.Score	University.Rating	SOP	LOR	CGPA	Research	Chance.of.Admit
1	337	118	4	4.5	4.5	9.65	1	0.92
2	324	107	4	4.0	4.5	8.87	1	0.76
3	316	104	3	3.0	3.5	8.00	1	0.72
4	322	110	3	3.5	2.5	8.67	1	0.80
5	314	103	2	2.0	3.0	8.21	0	0.65
9	302	102	1	2.0	1.5	8.00	0	0.50

	GRE.Score	TOEFL.Score	University.Rating	SOP	LOR	CGPA	Research	Chance.of.Admit
6	330	115	5	4.5	3.0	9.34	1	0.90
7	321	109	3	3.0	4.0	8.20	1	0.75
8	308	101	2	3.0	4.0	7.90	0	0.68
11	325	106	3	3.5	4.0	8.40	1	0.52
12	327	111	4	4.0	4.5	9.00	1	0.84
13	328	112	4	4.0	4.5	9.10	1	0.78

```
[1] "Train :  322"
```

```
[1] "Test :  78"
```

```

[23]: # First way of doing without any package
# linearMod <- lm(Chance.of.Admit ~ ., data=train_csv)
# summary(linearMod)

# Second way of doing using caret-lattice package
fitControl <- ml::trainControl(method = "repeatedcv", number = 3, repeats = 3)
regress <- ml::train(Chance.of.Admit ~ ., data = train_csv,
                     method = "lm", trControl = fitControl)
summary(regress)

```

```
Call:
lm(formula = .outcome ~ ., data = dat)

Residuals:
    Min       1Q   Median       3Q      Max
-0.25747 -0.02338  0.00914  0.03599  0.16426

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   -1.1679569   0.1329307  -8.786  < 2e-16 ***
GRE.Score       0.0013746   0.0006458   2.128  0.034080 *
TOEFL.Score     0.0027301   0.0012056   2.265  0.024222 *
University.Rating 0.0043070   0.0050164   0.859  0.391231
SOP            -0.0028466   0.0060599  -0.470  0.638860
LOR             0.0243216   0.0060538   4.018  7.36e-05 ***
CGPA           0.1233330   0.0134632   9.161  < 2e-16 ***
Research       0.0298928   0.0086033   3.475  0.000584 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06283 on 314 degrees of freedom
Multiple R-squared:  0.8084,    Adjusted R-squared:  0.8041
F-statistic: 189.3 on 7 and 314 DF,  p-value: < 2.2e-16
```

as we can see, the P-values indicates us that some of the data are quite more relevant than others as we predicted with visualization : * LOR and CGPA have a really low p-value which means that the confidence that they predict well is high. * GRE.Score, TOEFL.Score and Research are also usefull but with a less confidence (still really high in this case). * University.Rating, SOP have a high p-values meaning that these are not good variables (confidence) to determines the chances of admission

```
[24]: y_pred <- predict(regress, newdata = test_csv)
      y_true <- test_csv[,ncol(test_csv)]
```

```
[25]: print_res <- function(y_true, y_pred){
      print(paste0("RMSE => ", round(metrics::rmse(y_true, y_pred),3)))
      print(paste0("MSE  => ", round(metrics::mse(y_true, y_pred),3)))
      print(paste0("MBE  => ", round(metrics::bias(y_true, y_pred),3)))
      print(paste0("MAE  => ", round(mean(Metrics::ae(y_true, y_pred)),3)))
      print(paste0("MAPE => ", round(mean(Metrics::ape(y_true, y_pred)),3)))
      print(paste0("R2   => ", round(mlmetrics::R2_Score(y_true, y_pred),3)))
    }
```

```
[26]: print_res(y_true, y_pred)
```

```
[1] "RMSE => 0.068"
[1] "MSE  => 0.005"
[1] "MBE  => -0.005"
```

```
[1] "MAE  => 0.048"
[1] "MAPE => 0.077"
[1] "R2   => 0.718"
```

The result are quite good on the test set, meaning that all errors indicators are quite low.

Support Vector Machine (SVM)

```
[27]: fitControl <- ml::trainControl(method = "repeatedcv", number = 3, repeats = 3)
      svm <- ml::train(Chance.of.Admit ~ ., data = train_csv,
                      method = "svmLinear", trControl = fitControl)
      summary(svm)
```

```
Length Class  Mode
      1  ksvm   S4
```

```
[28]: y_pred <- predict(svm, newdata = test_csv)
      y_true <- test_csv[,ncol(test_csv)]
```

```
[29]: print_res(y_true, y_pred)
```

```
[1] "RMSE => 0.069"
[1] "MSE  => 0.005"
[1] "MBE  => -0.014"
[1] "MAE  => 0.047"
[1] "MAPE => 0.078"
[1] "R2   => 0.688"
```

The result are quite good on the test set, but not better than the basic linear regression.

XGBoost

```
[30]: fitControl <- ml::trainControl(method = "repeatedcv", number = 3, repeats = 3)
      svm <- ml::train(Chance.of.Admit ~ ., data = train_csv,
                      method = "xgbTree", trControl = fitControl)
      summary(svm)
```

	Length	Class	Mode
handle	1	xgb.Booster.handle	externalptr
raw	13359	-none-	raw
niter	1	-none-	numeric
call	5	-none-	call
params	8	-none-	list
callbacks	1	-none-	list
feature_names	7	-none-	character
nfeatures	1	-none-	numeric
xNames	7	-none-	character
problemType	1	-none-	character
tuneValue	7	data.frame	list
obsLevels	1	-none-	logical
param	0	-none-	list

```
[31]: y_pred <- predict(svm, newdata = test_csv)
      y_true <- test_csv[,ncol(test_csv)]
```

```
[32]: print_res(y_true, y_pred)
```

```
[1] "RMSE => 0.074"
[1] "MSE  => 0.005"
[1] "MBE  => -0.009"
[1] "MAE  => 0.052"
[1] "MAPE => 0.084"
[1] "R2   => 0.633"
```

1.4 Conclusion of the report

We had the opportunity to see both how to handle classification and regression models using Sklearn(Python) and caret(R). We also learnt about metrics and how to interpret them.